

How $224 \times 224 \times 3$ Image Transforms to 1280 Encoding in EfficientNet-B0

Overview

EfficientNet-B0 transforms a $224 \times 224 \times 3$ input image into a 1280-dimensional feature encoding through a series of carefully designed convolutional operations. This transformation happens through multiple stages, each progressively reducing spatial dimensions while increasing channel depth to capture increasingly complex features.

Complete Transformation Pipeline

Stage-by-Stage Breakdown

Stage 1: Initial Convolution

```
Input: 224x224x3 (Original RGB image)
Operation: Conv2D (3x3 kernel, stride=2, padding=1)
Output: 112x112x32
```

What happens:

- **Spatial reduction:** Image size halved ($224 \rightarrow 112$) due to stride=2
- **Channel expansion:** 3 RGB channels \rightarrow 32 feature channels
- **Feature learning:** Learns low-level features like edges and textures
- **Receptive field:** Each output pixel sees 3×3 region of input

Stage 2: MBConv1 Block

```
Input: 112x112x32
Operation: MBConv1 (expansion=1, kernel=3x3, 1 layer)
Output: 112x112x16
```

MBConv1 Details:

- **Expansion factor = 1:** No channel expansion ($32 \rightarrow 32$)
- **Depthwise Conv:** 3×3 depthwise convolution with 32 channels
- **Projection:** 1×1 convolution reduces channels ($32 \rightarrow 16$)
- **No residual:** Input/output channel mismatch prevents residual connection

Stage 3: MBConv6 Block ($\times 2$ layers)

Input: 112×112×16
Operation: MBConv6 (expansion=6, kernel=3×3, 2 layers)
Output: 112×112×24

MBConv6 Processing (per layer):

1. **Expansion:** 16 → 96 channels ($16 \times 6 = 96$)
2. **Depthwise:** 3×3 conv on 96 channels
3. **SE Block:** Squeeze-and-excitation attention
4. **Projection:** 96 → 24 channels

Stage 4: MBConv6 Block ($\times 2$ layers)

Input: 112×112×24
Operation: MBConv6 (expansion=6, kernel=5×5, 2 layers, stride=2)
Output: 56×56×40

Key Changes:

- **Spatial reduction:** 112×112 → 56×56 (stride=2 in first layer)
- **Larger kernels:** 5×5 for bigger receptive field
- **Channel growth:** 24 → 40 channels

Stage 5: MBConv6 Block ($\times 3$ layers)

Input: 56×56×40
Operation: MBConv6 (expansion=6, kernel=3×3, 3 layers, stride=2)
Output: 28×28×80

Processing:

- **Further spatial reduction:** 56×56 → 28×28
- **More layers:** 3 MBConv blocks for deeper feature learning
- **Channel doubling:** 40 → 80 channels

Stage 6: MBConv6 Block ($\times 3$ layers)

Input: 28×28×80
Operation: MBConv6 (expansion=6, kernel=5×5, 3 layers, stride=2)
Output: 14×14×112

Features:

- **Continued spatial reduction:** $28 \times 28 \rightarrow 14 \times 14$
- **Larger receptive field:** 5×5 kernels
- **High-level features:** Capturing complex patterns

Stage 7: MBConv6 Block ($\times 4$ layers)

```
Input: 14x14x112
Operation: MBConv6 (expansion=6, kernel=5x5, 4 layers)
Output: 14x14x192
```

Characteristics:

- **Same spatial size:** 14×14 maintained
- **Most layers:** 4 MBConv blocks for rich feature extraction
- **Channel growth:** $112 \rightarrow 192$ channels

Stage 8: MBConv6 Block ($\times 1$ layer)

```
Input: 14x14x192
Operation: MBConv6 (expansion=6, kernel=3x3, 1 layer, stride=2)
Output: 7x7x320
```

Final spatial reduction:

- **Smallest spatial size:** $14 \times 14 \rightarrow 7 \times 7$
- **Maximum channels:** 320 channels
- **High-level abstractions:** Complex semantic features

Stage 9: Final Convolution

```
Input: 7x7x320
Operation: Conv1x1 (1x1 kernel, 1280 filters)
Output: 7x7x1280
```

Purpose:

- **Channel expansion:** $320 \rightarrow 1280$ channels
- **Feature refinement:** Final feature processing
- **No spatial change:** Maintains 7×7 spatial size

Stage 10: Global Average Pooling

Input: $7 \times 7 \times 1280$
Operation: Global Average Pooling
Output: $1 \times 1 \times 1280 \rightarrow 1280$ (flattened)

Final transformation:

- **Spatial collapse:** $7 \times 7 \rightarrow 1 \times 1$ (average across spatial dimensions)
- **Feature vector:** Results in 1280-dimensional feature vector
- **Translation invariance:** Pooling provides spatial invariance

Detailed MBConv Block Analysis

MBConv6 Internal Structure (Example: $14 \times 14 \times 112 \rightarrow 14 \times 14 \times 192$)

1. Input Feature Map: $14 \times 14 \times 112$

2. Expansion Phase (1×1 Convolution):

- **Input:** $14 \times 14 \times 112$
- **Operation:** 1×1 Conv with 672 filters ($112 \times 6 = 672$)
- **Output:** $14 \times 14 \times 672$
- **Purpose:** Expand channels to create rich feature space
- **Activation:** Swish activation function

3. Depthwise Convolution:

- **Input:** $14 \times 14 \times 672$
- **Operation:** 5×5 depthwise convolution (672 separate 5×5 filters)
- **Output:** $14 \times 14 \times 672$
- **Purpose:** Spatial feature extraction per channel
- **Efficiency:** Much cheaper than standard convolution

4. Squeeze-and-Excitation (SE) Block:

- **Input:** $14 \times 14 \times 672$
- **Process:**
 - Global Average Pool: $14 \times 14 \times 672 \rightarrow 1 \times 1 \times 672$
 - FC layer 1: $672 \rightarrow 28$ (reduction ratio = 24)
 - ReLU activation
 - FC layer 2: $28 \rightarrow 672$
 - Sigmoid activation
 - Multiply with feature maps

- **Output:** $14 \times 14 \times 672$ (attention-weighted)
- **Purpose:** Channel attention mechanism

5. Projection Phase (1x1 Convolution):

- **Input:** $14 \times 14 \times 672$
- **Operation:** 1×1 Conv with 192 filters
- **Output:** $14 \times 14 \times 192$
- **Purpose:** Project back to desired output channels
- **No activation:** Linear projection to preserve information

6. Residual Connection (when applicable):

- **Condition:** Only when input/output dimensions match
- **Operation:** Element-wise addition of input and output
- **Purpose:** Gradient flow and feature reuse

Key Design Principles

1. Progressive Spatial Reduction

- $224 \times 224 \rightarrow 112 \times 112 \rightarrow 56 \times 56 \rightarrow 28 \times 28 \rightarrow 14 \times 14 \rightarrow 7 \times 7$
- Each reduction by factor of 2 (typical CNN pattern)
- Maintains computational efficiency

2. Channel Expansion

- $3 \rightarrow 32 \rightarrow 16 \rightarrow 24 \rightarrow 40 \rightarrow 80 \rightarrow 112 \rightarrow 192 \rightarrow 320 \rightarrow 1280$
- Generally increases with depth
- Captures increasingly complex features

3. Receptive Field Growth

- **Early stages:** Small kernels (3×3) for fine details
- **Middle stages:** Mixed kernels ($3 \times 3, 5 \times 5$) for medium features
- **Final stage:** Large effective receptive field for global context

4. Computational Efficiency

- **Depthwise separable convolutions:** Reduce parameters and computation
- **Inverted residuals:** Efficient information flow
- **Squeeze-and-excitation:** Lightweight attention mechanism

Mathematical Analysis

Total Parameter Reduction

MBCConv vs Standard Convolution parameter comparison:

- **Standard Conv:** $K \times K \times C_{in} \times C_{out}$ parameters
- **MBCConv:** $C_{in} \times E + K \times K \times (C_{in} \times E) + (C_{in} \times E) \times C_{out}$ parameters
- **Typical reduction:** 8-10x fewer parameters

Computational Complexity

For each MBCConv block:

1. **Expansion:** $H \times W \times C_{in} \times E$ multiplications
2. **Depthwise:** $H \times W \times K \times K \times (C_{in} \times E)$ multiplications
3. **Projection:** $H \times W \times (C_{in} \times E) \times C_{out}$ multiplications
4. **Total:** Much lower than equivalent standard convolution

Feature Map Memory

Progressive memory usage:

- **Stage 1:** $112 \times 112 \times 32 = 401,408$ values
- **Stage 4:** $56 \times 56 \times 40 = 125,440$ values
- **Stage 6:** $14 \times 14 \times 112 = 21,952$ values
- **Final:** $7 \times 7 \times 1280 = 62,720$ values

Why 1280 Dimensions?

1. Information Capacity

- **Rich representation:** 1280 dimensions can encode complex visual patterns
- **Balanced trade-off:** Not too high (overfitting) or too low (underfitting)
- **Transfer learning:** Good dimensionality for diverse downstream tasks

2. Architectural Design

- **NAS optimization:** Neural Architecture Search found this optimal
- **Compound scaling:** Fits well with EfficientNet's scaling principles
- **Hardware efficiency:** Good balance for modern GPUs

3. Empirical Performance

- **ImageNet accuracy:** Achieves excellent classification results
- **Transfer learning:** Works well across different domains
- **Medical imaging:** Particularly effective for chest X-rays and similar tasks

Feature Quality Characteristics

1. Hierarchical Features

- **Low-level:** Edges, textures, colors (early stages)
- **Mid-level:** Shapes, patterns, object parts (middle stages)
- **High-level:** Complex objects, semantic concepts (final stages)

2. Spatial Invariance

- **Global Average Pooling:** Provides translation invariance
- **Multi-scale processing:** Handles objects at different scales
- **Robust representations:** Less sensitive to small spatial variations

3. Channel Semantics

- **Specialized channels:** Different channels capture different features
- **SE attention:** Important channels get higher weights
- **Rich diversity:** 1280 channels provide comprehensive feature coverage

This transformation from $224 \times 224 \times 3$ to 1280 creates a powerful feature representation that captures both fine-grained details and high-level semantic information, making it ideal for medical image analysis and other computer vision tasks.